

Project Output Explanation


The output shows a visual comparison between the **input medical image**, the **ground truth (GT)** annotation, and the **model prediction** for a selected image slice.

- **Input:** The original brain MRI slice used as input to the model
- **GT (Ground Truth):** Manually annotated stroke lesion regions provided in the dataset
- **Prediction:** Segmentation mask generated by the trained **2D nnU-Net model**

The logs confirm that:

- All required files (image, GT, prediction) exist
- The shapes of the input image, ground truth, and prediction are identical (281, 352, 352), ensuring correct alignment and valid segmentation output

This indicates that the model successfully generated predictions and that the end-to-end pipeline (data loading → inference → visualization) is working correctly.



Starting visualization...

Data loaded

Slice index: 248

